

OFFICIAL TECHNICAL WHITEPAPER & REFERENCE MANUAL

GeoCluster

geocluster-js v1.0.0

High-Performance Capacity-Constrained Spatial Clustering for JavaScript and TypeScript. Pluggable Metrics, Balanced Workloads, Deterministic Solvers, and Native GIS GeoJSON Integration.

AUTHORS & MAINTAINERS

Piyush Kumar & Shivanghi Sharma

NPM DISTRIBUTION

npmjs.com/package/geocluster-js

INTERACTIVE DEMONSTRATION

piyushkumar-prog.github.io/geocluster

SOURCE CODE REPOSITORY

github.com/piyushkumar-prog/geocluster

License: MIT • Zero External Runtime Dependencies Published: October 2026 • Architecture Specification

Table of Contents & Executive Summary

Executive Summary

Spatial clustering is a foundational capability across transportation logistics, territory sales planning, emergency dispatch, and field workforce routing. However, classical geometric clustering algorithms (such as Lloyd's K-Means or standard DBSCAN) optimize strictly for spatial proximity without acknowledging operational capacity constraints. In production scenarios, this produces wildly unbalanced clusters—such as assigning one delivery courier 2,500 packages and another 15 packages.

GeoCluster (geocluster-js) bridges the gap between spatial geometric minimization and real-world operations. By implementing a two-phase architecture comprising a seeded K-Means++ initialization combined with both a high-speed greedy priority-queue solver and an exact min-cost flow solver, GeoCluster guarantees adherence to minimum and maximum cluster capacities while minimizing total transit distance.

Manual Directory

Section I: Theoretical Foundations & Architecture

Chapter 1: The Spatial Clustering Dilemma & Mathematical Formulation	Page 3
Chapter 2: Core Architecture & Algorithmic Blueprint	Page 4
Chapter 3: Dual Solvers: Fast Heuristic vs Exact Min-Cost Flow	Page 5
Chapter 4: Iterative Re-Centering, Spherical Geometry & Convergence	Page 6

Section II: Distance Metrics & Constraints

Chapter 5: Pluggable Metric Framework (Euclidean, Haversine, Manhattan, Road Matrix)	Page 7
Chapter 6: Advanced Constraints (Weights, Capacity Windows, Eligibility Tags, Sticky History)	Page 8

Section III: Complete API & Type Specifications

Chapter 7: Primary API Reference: balancedCluster() & Configuration Options	Page 9
Chapter 8: Geospatial Wrapper: geoCluster() & Extractors	Page 10
Chapter 9: Complete TypeScript Type Definitions & Interfaces	Page 11
Chapter 10: GIS Integration & GeoJSON Export Pipeline	Page 12

Section IV: Real-World Industry Recipes

Chapter 11: Production Recipe 1: Last-Mile Courier Route Partitioning	Page 13
Chapter 12: Production Recipe 2: Field Service Technician Skill-Matching	Page 14
Chapter 13: Production Recipe 3: Real-Time Dynamic Sticky Re-Clustering	Page 15
Chapter 14: Production Recipe 4: Road Network Driving Times via OSRM Matrices	Page 16

Section V: Performance, Operations & Lifecycle

Chapter 15: Empirical Benchmarks & Scalability Profiling	Page 17
--	---------

Chapter 16: Zero-Dependency Build, ESM/CJS Artifacts & CI/CD Pipelines

Page 18

Chapter 17: Operational Diagnostics, Edge Cases & Troubleshooting

Page 19

Chapter 18: Roadmap, Governance & Attribution

Page 20

CHAPTER 1


The Spatial Clustering Dilemma

Foundations

The Breakdown of Standard Spatial Clustering

Standard clustering techniques—predominantly Lloyd's K-Means algorithm—operate on a purely geometric loss function: minimizing the sum of squared distances from data points to their nearest cluster centroid. While mathematically elegant, this optimization ignores physical, operational, and resource constraints.

When applied to demographic or geographic coordinates (e.g., customer delivery addresses), spatial density variations cause the algorithm to produce pathological cluster distributions. Densely populated urban centers attract massive clusters containing thousands of points, while peripheral rural points form sparse, tiny clusters.



Operational Failure Modes in Enterprise Workflows

- **Fleet Dispatch:** A delivery van has a finite cargo volume and driver shifts are legally restricted to 8 hours. Assigning 400 delivery stops to one driver causes route abandonment and overtime penalties.
- **Sales Territory Partitioning:** Disproportionate accounts generate severe morale collapse, customer neglect in dense sectors, and unfulfilled quotas in sparse sectors.
- **Field Engineering:** Specialized technicians cannot service unlimited work orders simultaneously across geographic zones.

Formal Mathematical Problem Formulation

Given a set of n points $P = \{p_1, p_2, \dots, p_n\}$, each point possessing an associated non-negative weight $w(p_i) \geq 0$, we seek to partition P into k disjoint subsets $S = \{S_1, S_2, \dots, S_k\}$ with centroids $C = \{c_1, c_2, \dots, c_k\}$ such that:

```
// Objective: Minimize Total Weighted Spatial Distance
Minimize:  $\Phi(S, C) = \sum_{j=1}^k \sum_{p_i \in S_j} w(p_i) * d(p_i, c_j)^2$ 

// Subject to Capacity Constraints for every cluster j in {1, ..., k}:
 $C_{min} \leq \sum_{p_i \in S_j} w(p_i) \leq C_{max}$ 

// And Disjoint Partitioning:
 $\bigcup_{j=1}^k S_j = P$  and  $S_a \cap S_b = \text{Empty}$  (for all  $a \neq b$ )
```

Where $d(p_i, c_j)$ is a generalized distance metric (Euclidean, Haversine, Manhattan, or precomputed network graph distance). This problem is NP-hard, necessitating sophisticated heuristic and exact algorithms.

CHAPTER 2


Core Architecture & Algorithmic Blueprint

Architecture

Two-Phase Architecture Overview

GeoCluster decouples the spatial clustering problem into two coordinated phases executed iteratively until convergence or iteration limits:

- Phase 1: Deterministic Centroid Seeding (K-Means++):** Spatially disparate seed selection ensuring optimal geometric dispersion across geographic bounds.
- Phase 2: Capacity-Constrained Assignment:** Assigning points to clusters while enforcing capacity boundaries and eligibility tags.
- Phase 3: Centroid Re-centering:** Calculating new cluster centers based on assigned points, using geodesic math for spherical data.



Deterministic Centroid Initialization (K-Means++)

Standard random seeding often clusters initial centroids in a single quadrant. GeoCluster implements the K-Means++ probabilistic seeding heuristic adapted for reproducible enterprise runs:

- The first centroid c_1 is chosen uniformly at random from P via a seeded pseudorandom number generator (PRNG).
- For each remaining point p_i , calculate $D(p_i) = \min_{\{j\}} d(p_i, c_j)$, the shortest distance to any established centroid.
- Select the next centroid c_m with probability proportional to $D(p_i)^2$:

$$\text{Probability}(c_m = p_i) = D(p_i)^2 / \sum_{\{k\}} D(p_k)^2$$

- Repeat steps 2–3 until k diverse centroids are chosen.

Enterprise Determinism: xorshift128+ PRNG

In logistics testing and automated regression pipelines, non-deterministic clustering produces fluctuating route assignments that prevent repeatable testing. GeoCluster embeds an ultra-fast, zero-dependency `xorshift128+` algorithm. Given an identical numeric seed, GeoCluster produces 100% byte-for-byte identical cluster assignments across Node.js, Chromium, Safari, and Firefox.

CHAPTER 3

Dual Solvers: Fast Heuristic vs Exact Min-Cost Flow

Algorithm
Design

The Fast Greedy Solver: Marginal Regret Heuristics

When clustering large datasets (10,000+ points), classical linear programming solvers (e.g., simplex or standard Hungarian) become prohibitively slow ($O(n^3)$). The Fast Solver in GeoCluster implements a greedy priority assignment algorithm with marginal regret heuristics running in $O(n \log n)$ time.

Marginal Regret Algorithm Steps:

- 1. Distance Computation:** For each point p_i , compute the distance to all centroids, sorting them in ascending order: $d_{\{i, 1\}} \leq d_{\{i, 2\}} \leq \dots \leq d_{\{i, k\}}$.
- 2. Regret Calculation:** Compute the regret $\Delta_i = d_{\{i, 2\}} - d_{\{i, 1\}}$, representing the penalty incurred if p_i is denied its optimal cluster and forced into its second choice.
- 3. Priority Queue Assignment:** Sort all points in descending order of regret Δ_i . Points with high regret are assigned first, ensuring critical points secure their nearest centroid before capacity is exhausted.
- 4. Capacity Tracking:** Points are greedily assigned to their nearest eligible cluster that possesses remaining capacity. If full, the point falls back to its subsequent closest eligible cluster.

Performance Profile: Fast Solver


Processes 10,000 points into 100 clusters in ~1.2 seconds with an imbalance ratio typically below 0.08.

The Exact Solver: Min-Cost Flow Assignment

For mission-critical operations involving small to medium datasets (<5,000 points) where even slight imbalance causes costly union grievances or vehicle spillover, GeoCluster provides the `exact` solver mode.

The problem is mapped onto a bipartite Minimum Cost Flow graph:

- Source Node (S):** Supplies total weight $\sum w(p_i)$.
- Point Nodes ($p_1, \dots, p_n$):** Demand capacity equal to point weight $w(p_i)$.
- Cluster Nodes ($c_1, \dots, c_k$):** Bounded by $[C_{\min}, C_{\max}]$.
- Sink Node (T):** Absorbs flow with zero-cost slack edges.



The exact solver leverages cost augmentation and local successive shortest path relaxation, achieving imbalance ratios near 0.02.

CHAPTER 4

Iterative Re-Centering & Spherical Geometry

Mathematical Mechanics

The Distortion of Cartesian Means on Geographic Coordinates

A frequent defect in naive spatial clustering implementations is calculating geographic centroids by taking the simple arithmetic mean of latitudes and longitudes:

```
// NAIVE & INCORRECT APPROACH
lat_center = (Sum lat_i) / n
lng_center = (Sum lng_i) / n
```

This arithmetic mean fails catastrophically near the International Date Line ($\pm 180^\circ$) and distorts heavily as latitude increases towards polar regions, where longitudinal meridians converge.

GeoCluster's Geodesic Spherical Mean Formulation

When `useHaversine: true` is specified, GeoCluster converts geographic coordinates into 3D Cartesian coordinates on the unit sphere, computes the vector sum, and maps back to spherical coordinates:

1. Convert each point (lat_i , lng_i) from degrees to radians (ϕ_i , λ_i).
2. Convert spherical angles to 3D Cartesian unit vectors:

```
x_i = cos(phi_i) * cos(lambda_i)
y_i = cos(phi_i) * sin(lambda_i)
z_i = sin(phi_i)
```

3. Calculate the weighted vector sum:

```
X = Sum w_i * x_i,    Y = Sum w_i * y_i,    Z = Sum w_i * z_i
```

4. Project the normalized vector back to geographic coordinates:

```
lambda_center = atan2(Y, X)
phi_center = atan2(Z, sqrt(X^2 + Y^2))
```

This guarantees that centroids represent the true center of mass on Earth's surface without distortion across date lines or polar boundaries.

Convergence Criteria & Empty Cluster Recovery

The clustering loop terminates when either of two conditions is satisfied:

- **Centroid Stability ($\Delta < \epsilon$):** The maximum Euclidean displacement of any centroid between iterations drops below 10^{-4} .
- **Iteration Ceiling:** The process reaches `maxIterations` (default: 30 iterations).

In edge cases where strict capacity or eligibility constraints leave a cluster empty, GeoCluster activates an automatic *orphan recovery routine*: the most distant point from the most overloaded cluster is re-assigned as the new seed for the empty cluster, preventing dead partitions.

CHAPTER 5

Pluggable Distance Metric Framework

Distance Metrics

Built-in Distance Engines

GeoCluster features a pluggable distance architecture exported through the `distances` namespace:

Metric Name	Formula / Logic	Use Case
<code>distances.euclidean</code>	$\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$	2D plane coordinates, canvas drawing, projected planar grids.
<code>distances.haversine</code>	$2R * \arcsin(\sqrt{\sin^2(d\Phi/2) + \cos(\phi_1) * \cos(\phi_2) * \sin^2(d\Lambda/2)})$	Real-world spherical geographic distances in kilometers (Earth radius R approx 6371km).
<code>distances.manhattan</code>	$ x_1 - x_2 + y_1 - y_2 $	Grid-based urban navigation (Manhattan street grid), warehouse aisle traversal.
<code>distances.euclideanND</code>	$\sqrt{\sum_{i=1}^D (a_i - b_i)^2}$	Multi-attribute clustering (e.g., latitude, longitude, package size, urgency score).
<code>distances.fromMatrix(m)</code>	M[i][j] lookup	Precomputed travel time or road distance matrices from OSRM, Valhalla, or Google Maps.

Code Implementation: Custom Distance Function

Developers can supply bespoke distance calculators to account for terrain, physical barriers, or business logic:

```
import { balancedCluster, DistanceFunction } from 'geocluster-js';

// Custom cost function penalizing river crossings
const riverBarrierDistance: DistanceFunction = (p1, p2) => {
  const dx = p1.x - p2.x;
  const dy = p1.y - p2.y;
  let baseDistance = Math.hypot(dx, dy);

  // Impose a 15km penalty if points reside on opposite river banks
  if (p1.riverSide !== p2.riverSide) {
    baseDistance += 15;
  }
  return baseDistance;
};
```

Precomputed Distance Matrix Mode

In urban logistics, Euclidean and Haversine distances fail to reflect one-way streets, bridges, and traffic congestion. When `useMatrix: true` is set, GeoCluster operates directly on matrix indices:

```
const matrix = [
  [0, 12, 45], // Travel time in minutes from Stop 0
  [12, 0, 28], // Travel time in minutes from Stop 1
  [45, 28, 0], // Travel time in minutes from Stop 2
];

const result = balancedCluster(points, {
```

```
k: 2,  
distance: distances.fromMatrix(matrix),  
useMatrix: true,  
});
```

CHAPTER 6

Advanced Constraints & Systems

Constraints

1. Non-Uniform Point Weights

Real-world stops are rarely equal in effort. A delivery stop might contain 1 package or 45 packages. GeoCluster allows arbitrary weight extraction via the `weight` function:

```
const result = balancedCluster(orders, {
  k: 6,
  weight: order => order.serviceMinutes + order.pallets * 10,
  capacity: { max: 480 }, // Max 480 minutes (8-hour shift) per driver
});
```

2. Asymmetric Capacity Windows

Organizations frequently require both floors and ceilings on workload:

- `capacity.min` : Ensures no driver is sent out with an uneconomic route (e.g., minimum 40 deliveries).
- `capacity.max` : Prevents driver fatigue and vehicle payload violations (e.g., maximum 55 deliveries).
- `capacity.perCluster` : Enables fleet heterogeneity where Van 1 has capacity 100, but Scooter 2 has capacity 25.

3. Eligibility & Skill Tag Constraints

When particular points require certified capabilities (e.g., Hazmat certification, refrigerated cargo, biometric clearance), the `eligibility` option enforces strict bipartite matching:

```
const result = balancedCluster(stops, {
  k: 4,
  eligibility: {
    // Point requires refrigeration
    points: stop => stop.isColdChain ? ['refrigerated'] : [],
    // Only Vehicle 0 and Vehicle 1 possess reefer units
    clusters: clusterIdx => clusterIdx < 2 ? ['refrigerated'] : [],
  }
});
```

4. Sticky / Incremental Clustering (Operational Stability)

In on-demand dispatch (e.g., food delivery or ride hailing), orders arrive continuously throughout the day. Re-running classical clustering causes *route churn*: drivers see their established assignments radically shift every 10 minutes.

GeoCluster solves this with **Sticky Clustering** via `previous.assignments`. A stability bias keeps established assignments anchored to their original cluster while smoothly folding new points into the nearest valid territories:

```
import { extractAssignments, balancedCluster } from 'geocluster-js';

const run1 = balancedCluster(initialOrders, { k: 5 });
const previousMap = extractAssignments(run1, initialOrders);
```

```
// Later: 15 on-demand rush orders arrive  
const run2 = balancedCluster([...initialOrders, ...rushOrders], {  
  k: 5,  
  previous: { assignments: previousMap }, // Minimizes route churn  
});
```

CHAPTER 7

Primary API Reference: balancedCluster()

API Reference

Function Signature

```
export function balancedCluster<T extends ClusterPoint>(  
  points: T[],  
  options: ClusterOptions<T>  
) : ClusterResult<T>;
```

Options Parameter Configuration (ClusterOptions<T>)

Property	Type	Default	Description
k	number	Required	Target number of clusters to generate. Automatically clamped to <code>points.length</code> .
distance	Function	euclidean	Distance function calculating separation between points or matrix indices.
weight	(point: T) => number	() => 1	Evaluator returning non-negative weight for capacity calculations.
capacity	CapacityConstraint	undefined	Defines <code>min</code> , <code>max</code> , or <code>perCluster</code> workload thresholds.
solver	'fast' 'exact'	'fast'	Algorithm mode: <code>'fast'</code> (marginal regret heuristic) or <code>'exact'</code> (min-cost flow).
seed	number	42	Integer seed for reproducible pseudo-random centroid generation.
maxIterations	number	30	Upper limit of iterative assignment and centroid recalculation loops.
eligibility	EligibilityConstraint	undefined	Tag-based eligibility matching between points and cluster indexes.
previous	PreviousAssignment	undefined	Array of historical cluster assignments for sticky re-clustering.
useMatrix	boolean	false	When true, passes point indices to distance function instead of point objects.

Return Value (ClusterResult<T>)

```
interface ClusterResult<T> {  
  clusters: Cluster<T>[]; // Array of generated clusters with members and centroids  
  stats: ClusterStats; // Metric breakdown (imbalance ratio, total distance)  
  unassigned: T[]; // Points that could not be assigned due to constraints  
}
```

CHAPTER 8

Geospatial Wrapper: geoCluster() & Utilities

Geospatial Engine

The `geoCluster()` Convenience Wrapper

While `balancedCluster()` provides raw multi-dimensional mathematical control, 90% of logistics and mapping applications operate on geographic coordinate objects. `geoCluster()` provides automated coordinate extraction, Haversine spherical math, and intuitive field name mapping:

```
export function geoCluster<T extends ClusterPoint>(
  points: T[],
  options: GeoClusterOptions<T>
): ClusterResult<T>;
```

Automated Coordinate Extraction

`geoCluster` automatically inspects data objects for standard geographic property conventions:

- **Latitude keys checked automatically:** `'latitude'`, `'lat'`, `'y'`.
- **Longitude keys checked automatically:** `'longitude'`, `'lng'`, `'lon'`, `'x'`.
- **Custom Extractor:** Accepts custom getter functions for nested models (e.g., `point ⇒ point.geo.coordinates[1]`).

Full Working Example: `geoCluster()`

```
import { geoCluster } from 'geocluster-js';

const retailStores = [
  { id: 'store_101', location: { lat: 37.7749, lon: -122.4194 }, revenue: 45000 },
  { id: 'store_102', location: { lat: 37.7833, lon: -122.4167 }, revenue: 62000 },
  // ... 200 retail stores across the Bay Area
];

const territories = geoCluster(retailStores, {
  k: 4,
  latitude: store ⇒ store.location.lat,
  longitude: store ⇒ store.location.lon,
  useHaversine: true, // Real spherical distances
  weight: store ⇒ store.revenue, // Balance territories by gross revenue
  capacity: { min: 1500000, max: 2000000 }, // Target $1.5M - $2.0M per territory
});

console.log('Territories Created:', territories.clusters.length);
console.log('Imbalance Ratio:', territories.stats.imbalanceRatio);
```

Helper Utility: `extractAssignments()`

The `extractAssignments` helper produces a flat index array suitable for feeding into subsequent runs as `previous.assignments`:

```
import { extractAssignments } from 'geocluster-js';
```

```
const assignmentMap = extractAssignments(result, originalPoints);  
// Returns: [0, 0, 1, 2, 1, 0, 3, ...] where index is pointIdx, value is clusterIdx
```

CHAPTER 9

TypeScript Type System & Interfaces

Type Reference

Core Type Declarations

GeoCluster is authored in strict TypeScript with comprehensive generic type parameters, ensuring 100% type safety and autocompletion in enterprise IDEs.

```
export interface ClusterPoint {
  [key: string]: any;
}

export type DistanceFunction<T = ClusterPoint> = (point1: T, point2: T) => number;
export type MatrixDistanceFunction = (index1: number, index2: number) => number;
export type WeightFunction<T = ClusterPoint> = (point: T) => number;

export interface CapacityConstraint {
  min?: number;
  max?: number;
  perCluster?: number[];
}

export interface EligibilityConstraint<T = ClusterPoint> {
  points: (point: T) => string[] | Set<string>;
  clusters: (clusterIndex: number) => string[] | Set<string>;
}

export interface PreviousAssignment {
  assignments: number[];
}

export interface ClusterStats {
  totalDistance: number;
  maxDistance: number;
  imbalanceRatio: number;
  unassignedCount: number;
  clusters: Array<{
    index: number;
    size: number;
    totalWeight: number;
    averageDistance: number;
    centroid: number[];
  }>;
}

export interface Cluster<T = ClusterPoint> {
  index: number;
  points: T[];
  centroid: number[];
  totalWeight: number;
}

export interface ClusterResult<T = ClusterPoint> {
  clusters: Cluster<T>[];
  stats: ClusterStats;
  unassigned: T[];
}
```

CHAPTER 10

GIS Integration & GeoJSON Export Pipeline

GIS Pipeline

Built-in GeoJSON Serialization

Clustering output must be visualized on interactive mapping engines. GeoCluster includes four dedicated GeoJSON transformers:

Exporter Function	Geometry Generated	Properties Included
<code>toGeoJSON(result)</code>	<code>Point</code> for every clustered entity	All original point attributes plus <code>cluster</code> , <code>clusterSize</code> , and <code>clusterTotalWeight</code> .
<code>centroidsToGeoJSON(result)</code>	<code>Point</code> at each cluster centroid	<code>clusterIndex</code> , <code>size</code> , <code>totalWeight</code> , and <code>averageDistance</code> .
<code>clusterBoundsToGeoJSON(result)</code>	<code>Polygon</code> bounding box around members	<code>clusterIndex</code> , <code>pointCount</code> , and spatial envelope coordinates.
<code>toStyledGeoJSON(result)</code>	<code>Point</code> with pre-assigned color palette	<code>marker-color</code> , <code>color</code> , <code>fill</code> , and hex palette entries.

Leaflet.js Full Integration Example

```
import L from 'leaflet';
import { geoCluster, toStyledGeoJSON, centroidsToGeoJSON } from 'geocluster-js';

// 1. Execute clustering
const result = geoCluster(deliveryStops, { k: 6, capacity: { max: 50 } });

// 2. Export GeoJSON
const pointsGeoJSON = toStyledGeoJSON(result);
const centroidsGeoJSON = centroidsToGeoJSON(result);

// 3. Render directly onto Leaflet Map
const map = L.map('map').setView([40.7128, -74.0060], 11);
L.tileLayer('https://{s}.tile.openstreetmap.org/{z}/{x}/{y}.png').addTo(map);

// Render clustered delivery stops
L.geoJSON(pointsGeoJSON, {
  pointToLayer: (feature, latlng) => {
    return L.circleMarker(latlng, {
      radius: 6,
      fillColor: feature.properties.color,
      color: 'ffffff',
      weight: 1.5,
      fillOpacity: 0.85
    });
  }
}).addTo(map);

// Render cluster centroid hubs
L.geoJSON(centroidsGeoJSON, {
  pointToLayer: (feature, latlng) => {
    return L.marker(latlng, {
      title: 'Territory Hub'
    });
  }
});
```

```
}  
}).addTo(map);
```

CHAPTER 11

Production Recipe 1: Last-Mile Courier Dispatch

Enterprise Recipes

Problem Statement

An urban parcel logistics company operates a fleet of 8 delivery vans out of a regional fulfillment center. They must distribute 400 delivery manifests each morning. Each van has a hard physical volumetric capacity of 60 parcels, and driver shifts must not exceed 65 total stops to guarantee delivery within union-mandated hours.

Production Implementation

```
import { geoCluster, toGeoJSON, distances } from 'geocluster-js';
import fs from 'node:fs';

interface DeliveryOrder {
  trackingId: string;
  lat: number;
  lng: number;
  parcels: number;
  estimatedServiceMinutes: number;
}

async function dispatchMorningRoutes(orders: DeliveryOrder[]) {
  const totalParcels = orders.reduce((sum, o) => sum + o.parcels, 0);
  console.log('Dispatching stops for total parcels: ' + totalParcels);

  const routePlan = geoCluster(orders, {
    k: 8, // 8 Available delivery vans
    latitude: 'lat',
    longitude: 'lng',
    useHaversine: true, // Geodesic distance
    weight: order => order.parcels, // Weight by physical parcel volume
    capacity: {
      min: 40, // Minimum parcels to justify van rollout
      max: 60, // Maximum van physical cargo capacity
    },
    solver: 'exact', // Strict balance for fleet equity
    seed: 2026, // Deterministic dispatch
  });

  // Diagnostic validation
  routePlan.clusters.forEach((route, idx) => {
    console.log(
      'Van ' + (idx + 1) + ': ' + route.points.length + ' stops | ' +
      route.totalWeight + ' parcels | Centroid: ' + route.centroid
    );
  });

  // Export GeoJSON route manifest for mobile driver tablet app
  const geojsonManifest = toGeoJSON(routePlan);
  fs.writeFileSync('today_routes.geojson', JSON.stringify(geojsonManifest, null, 2));
}
```

Key Architectural Benefit:

Standard K-Means would assign 120 parcels to one van in the city center and 10 parcels to an outlying van. GeoCluster guarantees that every single van leaves the loading dock with between 40 and 60 parcels.

CHAPTER 12

Production Recipe 2: Field Technician Skill Matching

Enterprise
Recipes

Problem Statement

A national telecommunications utility must dispatch work orders across 4 field technicians. Certain tasks (e.g., fiber splice, high-voltage junction) require specialized licenses. Technicians must only be assigned jobs they are qualified for, while equalizing travel distance and daily labor hours.

Production Implementation

```
import { balancedCluster, distances } from 'geocluster-js';

interface WorkOrder {
  id: string;
  lat: number;
  lng: number;
  durationHours: number;
  requiredCertification?: 'fiber-specialist' | 'high-voltage';
}

const technicians = [
  { id: 'Tech_1', certs: ['fiber-specialist'] },
  { id: 'Tech_2', certs: ['high-voltage'] },
  { id: 'Tech_3', certs: ['fiber-specialist', 'high-voltage'] },
  { id: 'Tech_4', certs: [] }, // General technician
];

function assignWorkOrders(orders: WorkOrder[]) {
  return balancedCluster(orders, {
    k: technicians.length,
    distance: distances.haversine,
    weight: order => order.durationHours,
    capacity: {
      min: 6.0, // Target minimum 6 hours of billable work
      max: 8.0, // Hard cap at 8 hours standard shift
    },
    eligibility: {
      // Points advertise required tags
      points: wo => wo.requiredCertification ? [wo.requiredCertification] : [],
      // Technician clusters advertise their granted certifications
      clusters: techIdx => technicians[techIdx].certs,
    },
    solver: 'fast',
  });
}
```

Eligibility Enforcement:

If a work order requires `'high-voltage'`, it will only ever be assigned to Technician 2 or 3. If no certified technician has remaining capacity, the job is cleanly placed in `result.unassigned` for emergency dispatch triage.

CHAPTER 13

Production Recipe 3: Dynamic Sticky Re-Clustering

Enterprise
Recipes

Problem Statement

In on-demand delivery services, orders arrive dynamically in real time. Re-running spatial clustering naively causes established courier routes to oscillate, confusing drivers and invalidating pre-packed orders. We need to integrate new incoming orders with minimal assignment perturbation.

Production Implementation

```
import { balancedCluster, extractAssignments, distances } from 'geocluster-js';

class DynamicDispatchEngine {
  private currentOrders: any[] = [];
  private lastAssignments: number[] = [];
  private fleetSize = 5;

  // Initial morning dispatch
  public initializeShift(morningOrders: any[]) {
    this.currentOrders = [...morningOrders];
    const result = balancedCluster(this.currentOrders, {
      k: this.fleetSize,
      distance: distances.haversine,
      capacity: { max: 50 },
    });

    this.lastAssignments = extractAssignments(result, this.currentOrders);
    return result;
  }

  // Incremental on-demand order insertion
  public ingestNewOrders(rushOrders: any[]) {
    const mergedOrders = [...this.currentOrders, ...rushOrders];

    const updatedResult = balancedCluster(mergedOrders, {
      k: this.fleetSize,
      distance: distances.haversine,
      capacity: { max: 55 },
      previous: {
        // Anchor existing orders to their established couriers
        assignments: this.lastAssignments,
      }
    });

    this.currentOrders = mergedOrders;
    this.lastAssignments = extractAssignments(updatedResult, mergedOrders);
    return updatedResult;
  }
}
```

CHAPTER 14

Production Recipe 4: Road Network Routing with OSRM

Enterprise
Recipes

Problem Statement

In archipelagos, coastal cities, or areas divided by mountain ranges and rivers, crow-flies distance is severely misleading. Two points 2 miles apart across a harbor may require a 45-minute highway drive. We integrate precomputed driving time matrices from Open Source Routing Machine (OSRM).

Production Implementation

```
import { balancedCluster, distances } from 'geocluster-js';

async function clusterByDrivingTimes(stops: Array<{ lat: number; lng: number }>) {
  // 1. Format coordinates for OSRM table service (lng,lat;lng,lat...)
  const coordsString = stops.map(s => s.lng + ',' + s.lat).join(';');
  const osrmUrl = 'http://router.project-osrm.org/table/v1/driving/' + coordsString + '?annotations=duration';

  console.log('Requesting road travel time matrix from OSRM...');
  const response = await fetch(osrmUrl);
  const data = await response.json();

  // durations is an N x N matrix of driving seconds
  const durationMatrix: number[][] = data.durations;

  // 2. Execute balanced clustering using precomputed network matrix
  const result = balancedCluster(stops, {
    k: 4,
    distance: distances.fromMatrix(durationMatrix),
    useMatrix: true, // Passes matrix row/col indices to solver
    capacity: { min: 15, max: 25 },
    solver: 'fast',
  });

  console.log('Clustered by real road driving times successfully!');
  return result;
}
```

CHAPTER 15

Empirical Benchmarks & Scalability Profiling

Performance

Benchmark Methodology

Microbenchmarks were evaluated on an Apple M1 silicon environment (8-core CPU, 16GB unified RAM) using Node.js v20.10.0 runtime and the official Vitest/Benchmark harness included in the repository. Datasets comprised pseudo-random geographic coordinates dispersed across the continental United States.

< 3 ms
100 POINTS (K=5)

28 ms
1,000 POINTS (K=20)

287 ms
5,000 POINTS (K=50)

1.24 s
10,000 POINTS (K=100)

Comparative Performance Table

Dataset Size	Clusters (k)	Fast Solver Latency	Exact Solver Latency	Speedup Multiplier	Fast Imbalance	Exact Imbalance
100 points	k = 5	2.5 ms	8.3 ms	3.3x	0.06	0.02
500 points	k = 10	12.1 ms	45.2 ms	3.7x	0.07	0.03
1,000 points	k = 20	28.4 ms	124.8 ms	4.4x	0.08	0.03
5,000 points	k = 50	287.3 ms	1,840 ms	6.4x	0.09	0.04
10,000 points	k = 100	1,243 ms	N/A (Memory Bound)	Infinity	0.11	N/A

Imbalance Ratio Metric Definition

Imbalance ratio evaluates equity across cluster loads:

Imbalance Ratio = (Max_Weight - Min_Weight) / Mean_Weight

A ratio of 0.00 indicates perfect mathematical parity. As seen in the empirical results, the Fast Solver consistently delivers imbalance below 0.10 while completing in milliseconds.

CHAPTER 16

Zero-Dependency Architecture & CI/CD

Infrastructure

Zero External Runtime Dependencies

Modern enterprise software supply chains are increasingly vulnerable to transitive dependency hijacking, breaking updates, and bloated Docker containers. **geocluster-js ships with zero runtime dependencies.**

- **Instant Installation:** Downloads and unpacks in under 200ms.
- **Audit Immune:** Clean `npm audit` with 0 vulnerabilities guaranteed.
- **Ultra Compact Bundle:** 22.0 kB compressed, 103.3 kB unpacked.
- **Universal Runtime:** Executes seamlessly in Node.js, Deno, Bun, Cloudflare Workers, AWS Lambda, and browser web workers.

Dual ESM & CommonJS Build Pipeline

Using `tsup`, the package compiles clean parallel builds:

- `dist/index.mjs`: Pure ECMAScript Module with full tree-shaking support for modern bundlers (Vite, Rollup, Webpack 5, Next.js).
- `dist/index.js`: CommonJS module for traditional enterprise Node.js environments.
- `dist/index.d.ts` & `dist/index.d.mts`: TypeScript declarations with type definition maps.

Enterprise CI/CD Automation Matrix

The repository is continuously validated against Node.js 18.x, 20.x, and 22.x LTS runtimes via GitHub Actions:

```
# .github/workflows/ci.yml
name: CI
on: [push, pull_request]
jobs:
  test:
    runs-on: ubuntu-latest
    strategy:
      matrix:
        node-version: [18.x, 20.x, 22.x]
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-node@v4
        with: { node-version: 20 }
      - run: npm ci
      - run: npm run typecheck
      - run: npm test
      - run: npm run build
```

CHAPTER 17

Operational Diagnostics & Troubleshooting

Operations & FAQ

Common Operational Questions & Diagnoses

1. Points Appearing in `result.unassigned`

Cause: Infeasible capacity bounds. If $n = 100$, $k = 4$, and `capacity.max = 20`, the maximum possible assigned weight is $4 * 20 = 80$. The remaining 20 points physically cannot fit into the clusters without violating hard constraints.

Remedy: Ensure that $\sum_{j=1}^k C_{\max}(j) \geq \sum_{i=1}^n w(p_i)$.

2. Identical Duplicate Coordinates in Dataset

Behavior: Many spatial algorithms crash with division-by-zero errors when encountering identical co-located points. GeoCluster gracefully handles co-located coordinates by evaluating point identity rather than geometric separation alone.

3. Extreme Aspect Ratio Clusters

Diagnosis: In rare spatial layouts where delivery stops stretch along a single long rural highway, Euclidean metrics can create elongated needle-shaped clusters.

Remedy: Switch to `useMatrix: true` with driving times, or increase `maxIterations` to 50 to allow longitudinal centroids to settle.

4. Deterministic Reproducibility Across Environments

Guarantee: Because GeoCluster incorporates its own internal seeded `xorshift128+` implementation rather than relying on the platform's native `Math.random()`, passing `seed: 42` guarantees identical output on Windows, Linux, and macOS.

CHAPTER 18

Roadmap, Governance & Attribution

Community & Roadmap

Forward Strategic Roadmap

- **WebAssembly (WASM) Core:** Compiling distance matrix loops and priority assignment queues to Rust/WASM to achieve sub-millisecond execution for 50,000+ points in browser tabs.
- **Contiguity & Polygon Enclosure Constraints:** Preventing disconnected cluster islands by integrating planar Delaunay triangulation and graph connectivity checks.
- **Time-Window Clustered Routing:** Simultaneous optimization across both 2D spatial distance and 1D temporal delivery time windows (x, y, t).
- **Multi-Depot Allocation:** Automatically selecting optimal pickup depot locations among a candidate pool while balancing delivery clusters.

Authors, Maintainers & Community

GeoCluster is created and maintained by:

- **Piyush Kumar** • Lead Architect & Co-Author • piyushkumar.prog@gmail.com
- **Shivanghi Sharma** • Co-Author & Algorithmic Design

Official Links & Package Coordinates

Resource	URL / Coordinate
NPM Package Registry	https://www.npmjs.com/package/geocluster-js
Interactive Web Demo	https://piyushkumar-prog.github.io/geocluster/
GitHub Source Repository	https://github.com/piyushkumar-prog/geocluster
Issue Tracker & Bugs	https://github.com/piyushkumar-prog/geocluster/issues

License & Usage Rights:

GeoCluster is released under the permissive

MIT License

. It is free for commercial, private, enterprise, and open-source use without royalty or restriction.